

```
In [22]: !pip install sentencepiece
!pip install transformers
!pip install rich[jupyter]
```

Looking in indexes: <https://pypi.org/simple>, (<https://pypi.org/simple>,) <http://us-python.pkg.dev/colab-wheels/public/simple/> (<https://us-python.pkg.dev/colab-wheels/public/simple/>)
Requirement already satisfied: sentencepiece in /usr/local/lib/python3.10/dist-packages (0.1.99)
Looking in indexes: <https://pypi.org/simple>, (<https://pypi.org/simple>,) <http://us-python.pkg.dev/colab-wheels/public/simple/> (<https://us-python.pkg.dev/colab-wheels/public/simple/>)
Requirement already satisfied: transformers in /usr/local/lib/python3.10/dist-packages (4.29.0)
Requirement already satisfied: filelock in /usr/local/lib/python3.10/dist-packages (from transformers) (3.12.0)
Requirement already satisfied: huggingface-hub<1.0,>=0.11.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (0.14.1)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.10/dist-packages (from transformers) (1.22.4)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from transformers) (23.1)
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.10/dist-packages (from transformers) (6.0)

```
In [23]: import pandas as pd
df = pd.read_csv("/content/Conversation_Chatbot.csv")
```

```
In [24]: df.sample(10)
```

```
Out[24]:
```

	0	question	answer
83	84	What is the "Whiffenpoofs" at Yale University?	The "Whiffenpoofs" is the oldest collegiate a ...
84	85	The "Whiffenpoofs" is the oldest collegiate a ...	It's so nice
42	43	Yes, It's an interdisciplinary learning	Can, you elaborate more about the course?
108	109	Can you tell me for any grad program point of ...	It generally requires submitting transcripts, ...
20	21	it's not bad. there are a lot of people there.	good luck with that.
24	25	i'm absolutely lovely, thank you.	everything's been good with you?
5	6	i've been good. i'm in school right now.	what school do you go to?
1	2	i'm fine. how about yourself?	i'm pretty good. thanks for asking.
74	75	That's seems to be very tough	Yes, it is a highly reputed college.
126	127	It varies depending on the program.	Can you tell me for any grad program point of ...

```
In [25]: df["answer"] = "summarize: "+df["answer"]
```

In [26]: `df.head()`

Out[26]:

	0	question	answer
0	1	hi, how are you doing?	summarize: i'm fine. how about yourself?
1	2	i'm fine. how about yourself?	summarize: i'm pretty good. thanks for asking.
2	3	i'm pretty good. thanks for asking.	summarize: no problem. so how have you been?
3	4	no problem. so how have you been?	summarize: i've been great. what about you?
4	5	i've been great. what about you?	summarize: i've been good. i'm in school right...

```
In [27]: # Importing Libraries
import os
import numpy as np
import pandas as pd
import torch
import torch.nn.functional as F
from torch.utils.data import Dataset, DataLoader, RandomSampler, SequentialSampler
import os

# Importing the T5 modules from huggingface/transformers
from transformers import T5Tokenizer, T5ForConditionalGeneration

from rich.table import Table, Column
from rich import box
from rich.console import Console

# define a rich console logger
console=Console(record=True)

def display_df(df):
    """display dataframe in ASCII format"""

    console=Console()
    table = Table(Column("source_text", justify="center" ), Column("target_text"

    for i, row in enumerate(df.values.tolist()):
        table.add_row(row[0], row[1])

    console.print(table)

training_logger = Table(Column("Epoch", justify="center" ),
                        Column("Steps", justify="center"),
                        Column("Loss", justify="center"),
                        title="Training Status", pad_edge=False, box=box.ASCII)
```

```
In [28]: # Setting up the device for GPU usage
from torch import cuda
device = 'cuda' if cuda.is_available() else 'cpu'
```

```
In [29]: class YourDataSetClass(Dataset):
        """
        Creating a custom dataset for reading the dataset and
        loading it into the dataloader to pass it to the neural network for finetuni
        """

        def __init__(self, dataframe, tokenizer, source_len, target_len, source_text
            self.tokenizer = tokenizer
            self.data = dataframe
            self.source_len = source_len
            self.summ_len = target_len
            self.target_text = self.data[target_text]
            self.source_text = self.data[source_text]

        def __len__(self):
            return len(self.target_text)

        def __getitem__(self, index):
            source_text = str(self.source_text[index])
            target_text = str(self.target_text[index])

            #cleaning data so as to ensure data is in string type
            source_text = ' '.join(source_text.split())
            target_text = ' '.join(target_text.split())

            source = self.tokenizer.batch_encode_plus([source_text], max_length= self.
            target = self.tokenizer.batch_encode_plus([target_text], max_length= self.

            source_ids = source['input_ids'].squeeze()
            source_mask = source['attention_mask'].squeeze()
            target_ids = target['input_ids'].squeeze()
            target_mask = target['attention_mask'].squeeze()

            return {
                'source_ids': source_ids.to(dtype=torch.long),
                'source_mask': source_mask.to(dtype=torch.long),
                'target_ids': target_ids.to(dtype=torch.long),
                'target_ids_y': target_ids.to(dtype=torch.long)
            }
```

```
In [30]: def train(epoch, tokenizer, model, device, loader, optimizer):

    """
    Function to be called for training with the parameters passed from main func
    """

    model.train()
    for _,data in enumerate(loader, 0):
        y = data['target_ids'].to(device, dtype = torch.long)
        y_ids = y[:, :-1].contiguous()
        lm_labels = y[:, 1:].clone().detach()
        lm_labels[y[:, 1:] == tokenizer.pad_token_id] = -100
        ids = data['source_ids'].to(device, dtype = torch.long)
        mask = data['source_mask'].to(device, dtype = torch.long)

        outputs = model(input_ids = ids, attention_mask = mask, decoder_input_ids=
        loss = outputs[0]

        if _%10==0:
            training_logger.add_row(str(epoch), str(_), str(loss))
            console.print(training_logger)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
```

```
In [31]: def validate(epoch, tokenizer, model, device, loader):

    """
    Function to evaluate model for predictions

    """
    model.eval()
    predictions = []
    actuals = []
    with torch.no_grad():
        for _, data in enumerate(loader, 0):
            y = data['target_ids'].to(device, dtype = torch.long)
            ids = data['source_ids'].to(device, dtype = torch.long)
            mask = data['source_mask'].to(device, dtype = torch.long)

            generated_ids = model.generate(
                input_ids = ids,
                attention_mask = mask,
                max_length=150,
                num_beams=2,
                repetition_penalty=2.5,
                length_penalty=1.0,
                early_stopping=True
            )
            preds = [tokenizer.decode(g, skip_special_tokens=True, clean_up_token_spaces=True) for g in generated_ids]
            target = [tokenizer.decode(t, skip_special_tokens=True, clean_up_token_spaces=True) for t in y]
            if _%10==0:
                console.print(f'Completed {_}')

            predictions.extend(preds)
            actuals.extend(target)
    return predictions, actuals
```

In [31]:

```
In [32]: def T5Trainer(dataframe, source_text, target_text, model_params, output_dir="").

        """
        T5 trainer
        """

        # Set random seeds and deterministic pytorch for reproducibility
        torch.manual_seed(model_params["SEED"]) # pytorch random seed
        np.random.seed(model_params["SEED"]) # numpy random seed
        torch.backends.cudnn.deterministic = True

        # Logging
        console.log(f"[Model]: Loading {model_params['MODEL']}...\n")

        # tokenizer for encoding the text
        tokenizer = T5Tokenizer.from_pretrained(model_params["MODEL"])

        # Defining the model. We are using t5-base model and added a Language model
        # Further this model is sent to device (GPU/TPU) for using the hardware.
        model = T5ForConditionalGeneration.from_pretrained(model_params["MODEL"])
        model = model.to(device)

        # Logging
        console.log(f"[Data]: Reading data...\n")

        # Importing the raw dataset
        dataframe = dataframe[[source_text, target_text]]
        display_df(dataframe.head(2))

        # Creation of Dataset and Dataloader
        # Defining the train size. So 80% of the data will be used for training and
        train_size = 0.8
        train_dataset = dataframe.sample(frac=train_size, random_state = model_params["
        val_dataset = dataframe.drop(train_dataset.index).reset_index(drop=True)
        train_dataset = train_dataset.reset_index(drop=True)

        console.print(f"FULL Dataset: {dataframe.shape}")
        console.print(f"TRAIN Dataset: {train_dataset.shape}")
        console.print(f"TEST Dataset: {val_dataset.shape}\n")

        # Creating the Training and Validation dataset for further creation of Datal
        training_set = YourDataSetClass(train_dataset, tokenizer, model_params["MAX_
        val_set = YourDataSetClass(val_dataset, tokenizer, model_params["MAX_SOURCE

        # Defining the parameters for creation of dataloaders
        train_params = {
            'batch_size': model_params["TRAIN_BATCH_SIZE"],
            'shuffle': True,
            'num_workers': 0
        }
```

```

val_params = {
    'batch_size': model_params["VALID_BATCH_SIZE"],
    'shuffle': False,
    'num_workers': 0
}

# Creation of Dataloaders for testing and validation. This will be used down
training_loader = DataLoader(training_set, **train_params)
val_loader = DataLoader(val_set, **val_params)

# Defining the optimizer that will be used to tune the weights of the network
optimizer = torch.optim.Adam(params = model.parameters(), lr=model_params["

# Training Loop
console.log(f'[Initiating Fine Tuning]...\n')

for epoch in range(model_params["TRAIN_EPOCHS"]):
    train(epoch, tokenizer, model, device, training_loader, optimizer)

console.log(f"[Saving Model]...\n")
#Saving the model after training
path = os.path.join(output_dir, "model_files")
model.save_pretrained(path)
tokenizer.save_pretrained(path)

# evaluating test dataset
console.log(f"[Initiating Validation]...\n")
for epoch in range(model_params["VAL_EPOCHS"]):
    predictions, actuals = validate(epoch, tokenizer, model, device, val_loader)
    final_df = pd.DataFrame({'Generated Text':predictions,'Actual Text':actuals})
    final_df.to_csv(os.path.join(output_dir,'predictions.csv'))

console.save_text(os.path.join(output_dir,'logs.txt'))

console.log(f"[Validation Completed.]\n")
console.print(f""[Model] Model saved @ {os.path.join(output_dir, "model_files")}\n")
console.print(f""[Validation] Generation on Validation data saved @ {os.path.join(output_dir, "logs.txt")}\n")
console.print(f""[Logs] Logs saved @ {os.path.join(output_dir,'logs.txt')}\n")

```

```

In [33]: model_params={
    "MODEL": "t5-base",           # model_type: t5-base/t5-large
    "TRAIN_BATCH_SIZE": 8,       # training batch size
    "VALID_BATCH_SIZE": 8,       # validation batch size
    "TRAIN_EPOCHS": 3,           # number of training epochs
    "VAL_EPOCHS": 1,             # number of validation epochs
    "LEARNING_RATE": 1e-4,       # learning rate
    "MAX_SOURCE_TEXT_LENGTH": 512, # max length of source text
    "MAX_TARGET_TEXT_LENGTH": 50, # max length of target text
    "SEED": 42                   # set seed for reproducibility
}

```

```
In [34]: T5Trainer(dataframe=df[:127], source_text="question", target_text="answer", mo
```

```
[02:49:44] [Model]: Loading t5-base...
```

```
/usr/local/lib/python3.10/dist-packages/transformers/models/t5/tokenization_t5.py:163: FutureWarning: This tokenizer was incorrectly instantiated with a model max length of 512 which will be corrected in Transformers v5.
```

```
For now, this behavior is kept to avoid breaking backwards compatibility when padding/encoding with `truncation is True`.
```

```
- Be aware that you SHOULD NOT rely on t5-base automatically truncating your input to 512 when padding/encoding.
```

```
- If you want to encode/pad to sequences longer than 512 you can either instantiate this tokenizer with `model_max_length` or pass `max_length` when encoding/padding.
```

```
- To avoid this warning, please instantiate this tokenizer with `model_max_length` set to your preferred value.
```

```
warnings.warn(
```

```
[02:49:48] [Data]: Reading data...
```

Sample Data

source_text	target_text
hi, how are you doing?	summarize: i'm fine. how about yourself?
i'm fine. how about yourself?	summarize: i'm pretty good. thanks for asking

```
FULL Dataset: (127, 2)
```

```
TRAIN Dataset: (102, 2)
```

```
TEST Dataset: (25, 2)
```

```
[Initiating Fine Tuning]...
```

Training Status

Epoch	Steps	Loss
0	0	tensor(6.8346, device='cuda:0', grad_fn=<NllLossBackward0>)

Training Status

Epoch	Steps	Loss
0	0	tensor(6.8346, device='cuda:0', grad_fn=<NllLossBackward0>)
0	10	tensor(4.1578, device='cuda:0', grad_fn=<NllLossBackward0>)

Training Status

Epoch	Steps	Loss
0	0	tensor(6.8346, device='cuda:0', grad_fn=<NllLossBackward0>)

0	10	tensor(4.1578, device='cuda:0', grad_fn=<NllLossBackward0>)
1	0	tensor(4.4832, device='cuda:0', grad_fn=<NllLossBackward0>)

Training Status

Epoch	Steps	Loss
0	0	tensor(6.8346, device='cuda:0', grad_fn=<NllLossBackward0>)
0	10	tensor(4.1578, device='cuda:0', grad_fn=<NllLossBackward0>)
1	0	tensor(4.4832, device='cuda:0', grad_fn=<NllLossBackward0>)
1	10	tensor(3.5538, device='cuda:0', grad_fn=<NllLossBackward0>)

Training Status

Epoch	Steps	Loss
0	0	tensor(6.8346, device='cuda:0', grad_fn=<NllLossBackward0>)
0	10	tensor(4.1578, device='cuda:0', grad_fn=<NllLossBackward0>)
1	0	tensor(4.4832, device='cuda:0', grad_fn=<NllLossBackward0>)
1	10	tensor(3.5538, device='cuda:0', grad_fn=<NllLossBackward0>)
2	0	tensor(3.2713, device='cuda:0', grad_fn=<NllLossBackward0>)

Training Status

Epoch	Steps	Loss
0	0	tensor(6.8346, device='cuda:0', grad_fn=<NllLossBackward0>)
0	10	tensor(4.1578, device='cuda:0', grad_fn=<NllLossBackward0>)
1	0	tensor(4.4832, device='cuda:0', grad_fn=<NllLossBackward0>)
1	10	tensor(3.5538, device='cuda:0', grad_fn=<NllLossBackward0>)
2	0	tensor(3.2713, device='cuda:0', grad_fn=<NllLossBackward0>)
2	10	tensor(2.7085, device='cuda:0', grad_fn=<NllLossBackward0>)

[02:50:25] [Saving Model]...

[02:50:28] [Initiating Validation]...

Completed 0

[02:50:32] [Validation Completed.]

[Model] Model saved @ outputs/model_files

[Validation] Generation on Validation data saved @ outputs/predictions.csv

[Logs] Logs saved @ outputs/logs.txt

In [34]:

In [34]:

